

ECE 271A: Statistical Learning I Quiz Report

JunPyung Kim

PID: A69044427

University of California, San Diego

October 13, 2025

1 Quiz 1: Bayesian Classifier for Image Segmentation

1.1 Objective

The goal of this assignment is to implement a Bayesian classifier to segment an image of a cheetah from its background (grass). The classification is based on features derived from the Discrete Cosine Transform (DCT) of 8x8 image blocks.

1.2 Methodology and Results

1.2.1 Part (a): Prior Probabilities

The prior probabilities, $P(Y = \text{cheetah})$ and $P(Y = \text{grass})$, were estimated from the proportions of foreground (FG) and background (BG) samples in the provided training data, `TrainingSamplesDCT_8.mat`. The estimation is based on the formula:

$$P(Y = c) = \frac{N_c}{N_{\text{total}}}$$

where N_c is the number of samples for class c and N_{total} is the total number of training samples. Based on the 250 foreground samples and 1053 background samples, the calculated priors are as follows:

- $P(Y = \text{cheetah}) = \frac{250}{250+1053} \approx 0.1919$
- $P(Y = \text{grass}) = \frac{1053}{250+1053} \approx 0.8081$

1.2.2 Part (b): Class-Conditional Probabilities (Likelihoods)

The feature used for classification is the index (from 1 to 64, following a zig-zag scan) of the DCT coefficient with the second-largest absolute value. The class-conditional probabilities, $P(X|Y)$, were estimated by creating a normalized histogram of these features for each class. The resulting probability mass functions (PMFs) are shown in Figure 1. The plots show that for the cheetah class, the energy is concentrated in lower-frequency DCT coefficients (smaller indices), whereas for the grass class, the energy is more distributed towards higher-frequency coefficients.

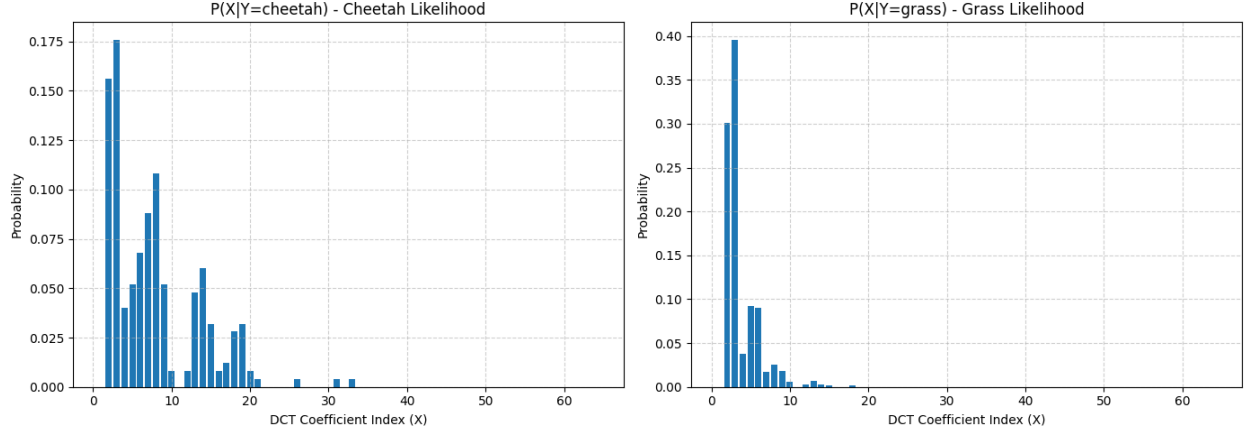


Figure 1: Estimated likelihoods $P(X|Y)$ for the Cheetah (left) and Grass (right) classes.

1.2.3 Part (c) & (d): Image Segmentation and Error Rate

The cheetah image was processed using a sliding 8x8 window. For each block, the feature was extracted and classified using the minimum probability of error rule: $\hat{Y} = \arg \max_Y P(X|Y)P(Y)$. The resulting segmentation mask is compared with the ground truth in Figure 2.

The probability of error was calculated by comparing the generated mask with the ground truth mask pixel by pixel. The formula used is:

$$P(\text{error}) = \frac{\text{Number of Mismatched Pixels}}{\text{Total Number of Pixels}}$$

The final computed probability of error, based on 11,566 mismatched pixels out of a total of 68,850, is: **16.80%**.

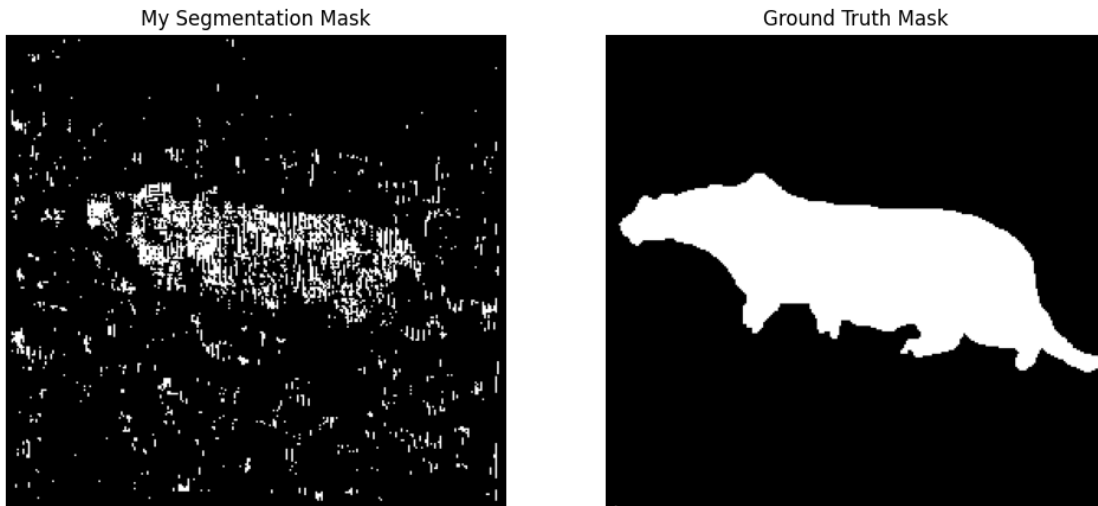


Figure 2: Left: Generated segmentation mask. Right: Ground truth mask.

1.3 Appendix: Source Code

As per the course guidelines, the source code used for this assignment is attached below.

```
1 import numpy as np
2 from scipy.io import loadmat
3 from scipy.fft import dctn
4 import imageio.v2 as imageio
5 import matplotlib.pyplot as plt
6
7 # --- 1. Data Loading ---
8 training_data = loadmat('TrainingSamplesDCT_8.mat')
9 TrainsampleDCT_FG = training_data['TrainsampleDCT_FG']
10 TrainsampleDCT_BG = training_data['TrainsampleDCT_BG']
11 zigzag_pattern = np.loadtxt('Zig-Zag Pattern.txt', dtype=int)
12 image = imageio.imread('cheetah.bmp', pilmode='L')
13 mask_true = imageio.imread('cheetah_mask.bmp')
14
15 # --- 2. Part (a): Estimate Prior Probabilities ---
16 n_fg = TrainsampleDCT_FG.shape[0]
17 n_bg = TrainsampleDCT_BG.shape[0]
18 n_total = n_fg + n_bg
19 prior_cheetah = n_fg / n_total
20 prior_grass = n_bg / n_total
21
22 print(f"Prior P(Y=cheetah): {prior_cheetah:.4f}")
23 print(f"Prior P(Y=grass): {prior_grass:.4f}")
24
25 # --- 3. Part (b): Estimate Class-Conditional Probabilities ---
26 def extract_features(data):
27     """Extracts the index of the 2nd largest absolute value for each sample."""
28     abs_data = np.abs(data)
29     sorted_indices = np.argsort(abs_data, axis=1)[: , :-1]
30     # Return the second column (index of 2nd largest value) + 1 for 1-based
    indexing
31     return sorted_indices[:, 1] + 1
32
33 features_fg = extract_features(TrainsampleDCT_FG)
34 features_bg = extract_features(TrainsampleDCT_BG)
35
36 bins = np.arange(1, 66)
37 hist_fg, _ = np.histogram(features_fg, bins=bins)
38 hist_bg, _ = np.histogram(features_bg, bins=bins)
39
40 likelihood_cheetah = hist_fg / n_fg
41 likelihood_grass = hist_bg / n_bg
42
43 # --- 4. Part (c): Classify the Image ---
44 image_float = image.astype(np.float64) / 255.0
45 height, width = image.shape
46 block_size = 8
47 segmentation_mask = np.zeros((height - block_size + 1, width - block_size + 1))
48
49 zigzag_flat = zigzag_pattern.flatten()
50 inverse_zigzag = np.empty_like(zigzag_flat)
51 inverse_zigzag[zigzag_flat] = np.arange(len(zigzag_flat))
52
53 for i in range(height - block_size + 1):
54     for j in range(width - block_size + 1):
```

```

55     block = image_float[i:i+block_size, j:j+block_size]
56     dct_block = dctn(block, type=2, norm='ortho')
57     dct_vector = dct_block.flatten()[inverse_zigzag]
58
59     feature_idx_0based = np.argsort(np.abs(dct_vector))[:, -1][1]
60     feature = feature_idx_0based + 1
61
62     # Apply Bayes Decision Rule using log-posteriors to avoid underflow
63     epsilon = 1e-10
64     posterior_cheetah = np.log(prior_cheetah) + np.log(likelihood_cheetah[
65     feature - 1] + epsilon)
66     posterior_grass = np.log(prior_grass) + np.log(likelihood_grass[feature -
67     1] + epsilon)
68
69     if posterior_cheetah > posterior_grass:
70         segmentation_mask[i, j] = 1 # 1 for cheetah
71     else:
72         segmentation_mask[i, j] = 0 # 0 for grass
73
74 # --- 5. Part (d): Compute Probability of Error ---
75 mask_true_binary = (mask_true / 255).astype(int)
76
77 mismatched_pixels = np.sum(segmentation_mask.astype(int) != mask_true_binary)
78 total_pixels = mask_true_binary.size
79 error_rate = mismatched_pixels / total_pixels
80
81 print(f"Total pixels: {total_pixels}")
82 print(f"Mismatched pixels: {mismatched_pixels}")
83 print(f"Probability of Error: {error_rate:.4f} or {error_rate * 100:.2f}%")

```

Listing 1: Python code for HW1