

# ECE 271A: Statistical Learning I Quiz Report

JunPyung Kim  
PID: A69044427  
University of California, San Diego

October 27, 2025

## 1 Quiz 2: Gaussian Bayesian Classifier for Image Segmentation

### 1.1 Objective

The goal of this assignment is to extend the previous Bayesian classifier by modeling the class-conditional densities,  $P_{X|Y}(x|c)$ , as multivariate Gaussian distributions. We will compare the performance of a full 64-dimensional Gaussian model against a reduced 8-dimensional model built using the most discriminative features.

### 1.2 Methodology and Results

#### 1.2.1 Part (a): Prior Probabilities

The prior probabilities,  $P(Y = \text{cheetah})$  and  $P(Y = \text{grass})$ , were re-evaluated using the new `TrainingSamplesDCT_8_new.mat` data. The Maximum Likelihood Estimate (MLE) for the parameter  $p_c$  of a categorical distribution is given by:

$$\hat{p}_c = P(Y = c) = \frac{N_c}{N_{\text{total}}}$$

where  $N_c$  is the number of samples for class  $c$ . Based on the 250 foreground (cheetah) and 1053 background (grass) samples:

- $P(Y = \text{cheetah}) = \frac{250}{250+1053} = \frac{250}{1303} \approx 0.1919$
- $P(Y = \text{grass}) = \frac{1053}{250+1053} = \frac{1053}{1303} \approx 0.8081$

These results are identical to the priors computed in Quiz 1.

#### 1.2.2 Part (b): Class-Conditional Parameters and Feature Selection

The class-conditional densities were modeled as 64-dimensional Gaussian distributions,  $P_{X|Y}(x|c) \sim \mathcal{N}(\boldsymbol{\mu}_c, \boldsymbol{\Sigma}_c)$ . The MLE parameters for the mean vector  $\boldsymbol{\mu}_c$  and covariance matrix  $\boldsymbol{\Sigma}_c$  were computed from the training data. The plot for all 64 marginal densities is shown in Figure 1.

Feature selection was performed quantitatively using the \*\*Symmetric Kullback-Leibler (KL) Divergence\*\*. This metric measures the "distance" between the 1D marginal distributions  $P(X_k|\text{cheetah})$  and  $P(X_k|\text{grass})$  for each feature. The 8 features with the highest KL divergence (most separable) and the 8 with the lowest (least separable) were identified.




Figure 1: Marginal densities  $P(X_k|Y)$  for all 64 DCT features.

- **Best 8 Features (by KL Div):** [0, 31, 26, 24, 39, 32, 17, 40]
- **Worst 8 Features (by KL Div):** [63, 62, 58, 2, 4, 61, 3, 59]

The marginal densities for these selected features are plotted in Figure 2.

### 1.2.3 Part (c): Image Segmentation and Performance

The `cheetah.bmp` image was classified using a sliding 8x8 window. At each pixel, the 64-D DCT vector for the corresponding block was extracted and classified using the Bayes Decision Rule for




Figure 2: Marginal densities  $P(X_k|Y)$  for the 8 best (top) and 8 worst (bottom) features, as selected by Symmetric KL Divergence.

minimum error:

$$\hat{Y} = \arg \max_{Y \in \{\text{cheetah, grass}\}} [\log P_{X|Y}(x|Y) + \log P(Y)]$$

This classification was performed twice:

1. Using the full 64-dimensional Gaussian models.
2. Using 8-dimensional Gaussian models built from the best 8 features.

The resulting probability of error for each classifier, computed by comparing the output mask to the ground truth, was:

- **64-D Classifier Error:** 0.1450 (14.50%)
- **8-D Classifier Error:** 0.0748 (7.48%)

#### 1.2.4 Discussion: Explaining the Results

As shown by the error rates and the segmentation masks (Figures 3 and 4), the **8-dimensional classifier performs significantly better** than the 64-dimensional one.

This result is a classic example of the “Curse of Dimensionality.” With a limited number of training samples (250 for cheetah, 1053 for grass), accurately estimating the parameters for a 64-dimensional Gaussian distribution is extremely difficult. The 64x64 covariance matrix has  $\frac{64 \times 65}{2} = 2080$  unique parameters. Estimating this many parameters from only 250 foreground samples leads to a model that is highly overfit to the training data and does not generalize well.

The 8-dimensional model, in contrast, requires estimating an 8x8 covariance matrix (36 parameters). This model is much simpler, its parameters can be estimated more robustly from the available data, and it consequently generalizes better to the unseen test image, resulting in a lower probability of error.




Figure 3: Segmentation using the full 64-D Gaussian model. Left: Predicted Mask. Right: Ground Truth. (Error: 14.50%)




Figure 4: Segmentation using the 8-D Gaussian model (best features). Left: Predicted Mask. Right: Ground Truth. This model achieves a lower error rate. (Error: 7.48%)

### 1.3 Appendix: Source Code

The Python source code used for this assignment is attached below.

```

1 import numpy as np
2 import scipy.io
3 import scipy.stats

```

```

4 import imageio.v3 as imageio
5 import matplotlib
6 import os
7 import warnings
8 from scipy.fftpack import dctn
9 from tqdm import tqdm
10
11 try:
12     matplotlib.use('Agg')
13 except ImportError:
14     pass
15
16 import matplotlib.pyplot as plt
17
18 warnings.filterwarnings('ignore')
19 np.random.seed(42)
20
21 BLOCK_SIZE = 8
22 DATA_FILE = 'data/TrainingSamplesDCT_8_new.mat'
23 ZIG_ZAG_FILE = 'data/Zig-Zag Pattern.txt'
24 IMAGE_FILE = 'data/cheetah.bmp'
25 MASK_FILE = 'data/cheetah_mask.bmp'
26 OUTPUT_DIR = 'hw2/output/'
27
28
29 def load_zig_zag_map(filepath):
30     """Loads the zig-zag scan pattern."""
31     zig_zag_pattern = np.loadtxt(filepath, dtype=int)
32     return np.argsort(zig_zag_pattern.flatten())
33
34
35 def dct2(block):
36     """Compute 2D DCT of an 8x8 block."""
37     return dctn(block, type=2, norm='ortho')
38
39
40 def compute_mle_parameters(data):
41     """
42         Compute MLE for mean and covariance matrix (ddof=0).
43     """
44     mean = np.mean(data, axis=0)
45     cov = np.cov(data, rowvar=False, ddof=0)
46     reg = np.eye(cov.shape[0]) * 1e-6
47     cov_reg = cov + reg
48     return mean, cov_reg
49
50
51 def compute_kl_divergence(p, q):
52     """Compute symmetric KL divergence between two 1D Gaussian distributions."""
53     mu_p, var_p = p
54     mu_q, var_q = q
55     epsilon = 1e-10
56     var_p = max(var_p, epsilon)
57     var_q = max(var_q, epsilon)
58
59     kl_pq = 0.5 * (np.log(var_q / var_p) + (var_p + (mu_p - mu_q) ** 2) / var_q - 1)
60     kl_qp = 0.5 * (np.log(var_p / var_q) + (var_q + (mu_q - mu_p) ** 2) / var_p - 1)

```

```

61     return kl_pq + kl_qp
62
63
64 def select_features_kl(fg_data, bg_data):
65     """Select best and worst 8 features based on symmetric KL divergence."""
66     kl_divergences = []
67     for i in range(64):
68         mean_fg = np.mean(fg_data[:, i])
69         var_fg = np.var(fg_data[:, i], ddof=0)
70         mean_bg = np.mean(bg_data[:, i])
71         var_bg = np.var(bg_data[:, i], ddof=0)
72         kl = compute_kl_divergence((mean_fg, var_fg), (mean_bg, var_bg))
73         kl_divergences.append(kl)
74
75     sorted_indices = np.argsort(kl_divergences)
76     worst_8_indices = sorted_indices[:8]
77     best_8_indices = sorted_indices[-8:][::-1] # Top 8
78     return best_8_indices, worst_8_indices
79
80
81 def plot_best_worst_features(fg_data, bg_data, best_indices, worst_indices,
82                             output_path):
83     """Plot marginals for best and worst 8 features."""
84     fig, axes = plt.subplots(2, 8, figsize=(20, 7))
85     fig.suptitle("Marginal Densities (Best and Worst 8 Features by KL Div.)",
86                  fontsize=16)
86
87     for i, idx in enumerate(best_indices):
88         ax = axes[0, i]
89         mean_fg = np.mean(fg_data[:, idx])
90         std_fg = np.std(fg_data[:, idx], ddof=0)
91         mean_bg = np.mean(bg_data[:, idx])
92         std_bg = np.std(bg_data[:, idx], ddof=0)
93
93         x_min = min(mean_fg - 4 * std_fg, mean_bg - 4 * std_bg)
94         x_max = max(mean_fg + 4 * std_fg, mean_bg + 4 * std_bg)
95         x = np.linspace(x_min, x_max, 200)
96
97         ax.plot(x, scipy.stats.norm.pdf(x, mean_fg, std_fg), 'b-', label='Cheetah')
98
99     ax.plot(x, scipy.stats.norm.pdf(x, mean_bg, std_bg), 'r-', label='Grass')
100    ax.set_title(f'Best Feature {idx}', fontsize=10)
101    if i == 0: ax.legend()
102
103    for i, idx in enumerate(worst_indices):
104        ax = axes[1, i]
105        mean_fg = np.mean(fg_data[:, idx])
106        std_fg = np.std(fg_data[:, idx], ddof=0)
107        mean_bg = np.mean(bg_data[:, idx])
108        std_bg = np.std(bg_data[:, idx], ddof=0)
109
110        x_min = min(mean_fg - 4 * std_fg, mean_bg - 4 * std_bg)
111        x_max = max(mean_fg + 4 * std_fg, mean_bg + 4 * std_bg)
112        x = np.linspace(x_min, x_max, 200)
113
113        ax.plot(x, scipy.stats.norm.pdf(x, mean_fg, std_fg), 'b-')
114        ax.plot(x, scipy.stats.norm.pdf(x, mean_bg, std_bg), 'r-')
115        ax.set_title(f'Worst Feature {idx}', fontsize=10)
116

```

```

117 plt.tight_layout(rect=[0, 0.03, 1, 0.95])
118 plt.savefig(output_path, dpi=150)
119 plt.close(fig)
120
121
122 def plot_all_features(fg_data, bg_data, output_path):
123     """Plot marginals for all 64 features."""
124     fig, axes = plt.subplots(8, 8, figsize=(22, 22))
125     fig.suptitle("Marginal Densities for All 64 Features", fontsize=16)
126
127     for i in range(64):
128         ax = axes[i // 8, i % 8]
129
130         mean_fg = np.mean(fg_data[:, i])
131         std_fg = np.std(fg_data[:, i], ddof=0)
132         mean_bg = np.mean(bg_data[:, i])
133         std_bg = np.std(bg_data[:, i], ddof=0)
134
135         std_fg = max(std_fg, 1e-6)
136         std_bg = max(std_bg, 1e-6)
137
138         x_min = min(mean_fg - 4 * std_fg, mean_bg - 4 * std_bg)
139         x_max = max(mean_fg + 4 * std_fg, mean_bg + 4 * std_bg)
140         x = np.linspace(x_min, x_max, 100)
141
142         ax.plot(x, scipy.stats.norm.pdf(x, mean_fg, std_fg), 'b-', label='Cheetah')
143         if i == 0 else "")
144         ax.plot(x, scipy.stats.norm.pdf(x, mean_bg, std_bg), 'r-', label='Grass')
145         if i == 0 else "")
146         ax.set_title(f'Feature {i}', fontsize=10)
147         ax.tick_params(labelsize=8)
148
149     fig.legend(loc='upper right')
150     plt.tight_layout(rect=[0, 0.03, 1, 0.95])
151     plt.savefig(output_path, dpi=150)
152     plt.close(fig)
153
154
155 def plot_segmentation_results(predicted_mask, true_mask, title, output_path):
156     """Plot segmentation results."""
157     fig, axes = plt.subplots(1, 2, figsize=(12, 6))
158     axes[0].imshow(predicted_mask, cmap='gray')
159     axes[0].set_title('Predicted Mask', fontsize=14)
160     axes[0].axis('off')
161     axes[1].imshow(true_mask, cmap='gray')
162     axes[1].set_title('Ground Truth Mask', fontsize=14)
163     axes[1].axis('off')
164     plt.suptitle(title, fontsize=16)
165     plt.tight_layout(rect=[0, 0.03, 1, 0.95])
166     plt.savefig(output_path, dpi=150)
167     plt.close(fig)
168
169
170 def extract_dct_vectors_sliding_window(image, zig_zag_map):
171     """
172     Processes an image using a sliding window.
173     Returns a (H*W, 64) array of DCT vectors.
174     """
175
176     img = np.float32(image)

```

```

174     h, w = img.shape
175
176     img_padded = np.pad(img, ((0, BLOCK_SIZE - 1), (0, BLOCK_SIZE - 1)), mode='reflect')
177     all_dct_vectors = np.zeros((h * w, BLOCK_SIZE * BLOCK_SIZE))
178
179     idx = 0
180     for i in tqdm(range(h), desc="Extracting DCT"):
181         for j in range(w):
182             block = img_padded[i:i + BLOCK_SIZE, j:j + BLOCK_SIZE]
183             block_dct = dct2(block)
184             all_dct_vectors[idx] = block_dct.flatten()[zig_zag_map]
185             idx += 1
186
187     return all_dct_vectors, h, w
188
189
190 def classify_blocks_vectorized(X_test, mean_fg, cov_fg, mean_bg, cov_bg,
191     log_prior_fg, log_prior_bg):
191     """
192     Classifies a set of test vectors using Bayesian decision rule.
193     """
194     log_ll_fg = scipy.stats.multivariate_normal.logpdf(X_test, mean=mean_fg, cov=cov_fg, allow_singular=True)
195     log_ll_bg = scipy.stats.multivariate_normal.logpdf(X_test, mean=mean_bg, cov=cov_bg, allow_singular=True)
196
197     g_fg = log_ll_fg + log_prior_fg
198     g_bg = log_ll_bg + log_prior_grass
199     decisions = (g_fg > g_bg).astype(int)
200     return decisions
201
202
203 def compute_error_rate(predicted_mask, true_mask, prior_fg, prior_bg):
204     """
205     Compute Bayesian probability of error (weighted).
206     P(error) = P(error|cheetah)P(cheetah) + P(error|grass)P(grass)
207     """
208     predicted_binary = (predicted_mask > 0.5).astype(int)
209     true_binary = (true_mask > 0.5).astype(int)
210
211     fg_pixels_total = np.sum(true_binary == 1)
212     bg_pixels_total = np.sum(true_binary == 0)
213
214     p_error_given_fg = 0.0
215     if fg_pixels_total > 0:
216         fg_misclassified = np.sum((true_binary == 1) & (predicted_binary == 0))
217         p_error_given_fg = fg_misclassified / fg_pixels_total
218
219     p_error_given_bg = 0.0
220     if bg_pixels_total > 0:
221         bg_misclassified = np.sum((true_binary == 0) & (predicted_binary == 1))
222         p_error_given_bg = bg_misclassified / bg_pixels_total
223
224     total_error = (p_error_given_fg * prior_fg) + (p_error_given_bg * prior_bg)
225     return total_error
226
227
228 def main():

```

```

229     os.makedirs(OUTPUT_DIR, exist_ok=True)
230
231     zig_zag_map = load_zig_zag_map(ZIG_ZAG_FILE)
232     train_data = scipy.io.loadmat(DATA_FILE)
233     fg_data = train_data['TrainsampleDCT_FG']
234     bg_data = train_data['TrainsampleDCT_BG']
235
236     image = imageio.imread(IMAGE_FILE, mode='L')
237     true_mask = imageio.imread(MASK_FILE)
238     true_mask = (true_mask > 127).astype(int)
239
240     n_fg = fg_data.shape[0]
241     n_bg = bg_data.shape[0]
242     n_total = n_fg + n_bg
243     prior_cheetah = n_fg / n_total
244     prior_grass = n_bg / n_total
245     log_prior_cheetah = np.log(prior_cheetah)
246     log_prior_grass = np.log(prior_grass)
247
248     print("\n" + "=" * 70)
249     print("Problem 6(a): Prior Probabilities")
250     print(f"P(Y=cheetah): {prior_cheetah:.6f} (N={n_fg})")
251     print(f"P(Y=grass): {prior_grass:.6f} (N={n_bg})")
252
253     print("\n" + "=" * 70)
254     print("Problem 6(b): ML Parameters and Feature Selection")
255
256     mean_fg_64, cov_fg_64 = compute_mle_parameters(fg_data)
257     mean_bg_64, cov_bg_64 = compute_mle_parameters(bg_data)
258
259     best_8_indices, worst_8_indices = select_features_kl(fg_data, bg_data)
260
261     print(f"Best 8 features (KL Div): {best_8_indices.tolist()}")
262     print(f"Worst 8 features (KL Div): {worst_8_indices.tolist()}")
263
264     plot_features_filename = os.path.join(OUTPUT_DIR, 'best_worst_8_features_KL.png')
265     plot_best_worst_features(fg_data, bg_data, best_8_indices, worst_8_indices,
266                               plot_features_filename)
267     print(f"Feature plots saved to {plot_features_filename}")
268
269     plot_all_filename = os.path.join(OUTPUT_DIR, 'marginal_densities_all_64.png')
270     plot_all_features(fg_data, bg_data, plot_all_filename)
271     print(f"All 64 feature plots saved to {plot_all_filename}")
272
273
274     print("\n" + "=" * 70)
275     print("Problem 6(c): Classification")
276     X_test_64, img_h, img_w = extract_dct_vectors_sliding_window(image,
277                                                                     zig_zag_map)
278
279     print("Classifying with 64D Gaussians...")
280     decisions_64_flat = classify_blocks_vectorized(X_test_64,
281                                                    mean_fg_64, cov_fg_64,
282                                                    mean_bg_64, cov_bg_64,
283                                                    log_prior_cheetah,
284                                                    log_prior_grass)
285     mask_64d = decisions_64_flat.reshape(img_h, img_w)

```

```

284     error_64d = compute_error_rate(mask_64d, true_mask, prior_cheetah, prior_grass
285 )
286
286 plot_64d_filename = os.path.join(OUTPUT_DIR, 'segmentation_64d_KL.png')
287 plot_segmentation_results(mask_64d, true_mask,
288                             f'64D Gaussian (KL) (Error: {error_64d:.4f})',
289                             plot_64d_filename)
290 print(f"64D plot saved to {plot_64d_filename}")
291
292 print("Classifying with 8D Gaussians (best features)...")
293 X_test_8 = X_test_64[:, best_8_indices]
294
295 mean_fg_8 = mean_fg_64[best_8_indices]
296 mean_bg_8 = mean_bg_64[best_8_indices]
297 cov_fg_8 = cov_fg_64[np.ix_(best_8_indices, best_8_indices)]
298 cov_bg_8 = cov_bg_64[np.ix_(best_8_indices, best_8_indices)]
299
300 decisions_8_flat = classify_blocks_vectorized(X_test_8,
301                                               mean_fg_8, cov_fg_8,
302                                               mean_bg_8, cov_bg_8,
303                                               log_prior_cheetah,
304                                               log_prior_grass)
304 mask_8d = decisions_8_flat.reshape(img_h, img_w)
305 error_8d = compute_error_rate(mask_8d, true_mask, prior_cheetah, prior_grass)
306
307 plot_8d_filename = os.path.join(OUTPUT_DIR, 'segmentation_8d_KL.png')
308 plot_segmentation_results(mask_8d, true_mask,
309                             f'8D Gaussian (KL) (Error: {error_8d:.4f})',
310                             plot_8d_filename)
311 print(f"8D plot saved to {plot_8d_filename}")
312
313 print("\n" + "=" * 70)
314 print("Final Explanation of Results")
315 print("=" * 70)
316 print(f"64D Classifier Bayesian Error: {error_64d:.6f}")
317 print(f" 8D Classifier Bayesian Error: {error_8d:.6f}")
318
319 if error_8d < error_64d:
320     print("\nSUCCESS!")
321 else:
322     print("\nSomething went wrong!")
323
324 if __name__ == "__main__":
325     main()

```

Listing 1: Python code for HW2 (hw2/hw2\_solution.py)