

ECE 271A: Statistical Learning I Quiz Report

JunPyung Kim
PID: A69044427
University of California, San Diego

October 27, 2025

1 Quiz 2: Gaussian Bayesian Classifier for Image Segmentation

1.1 Objective

The goal of this assignment is to extend the previous Bayesian classifier by modeling the class-conditional densities, $P_{X|Y}(x|c)$, as multivariate Gaussian distributions. We will compare the performance of a full 64-dimensional Gaussian model against a reduced 8-dimensional model built using the most discriminative features.

1.2 Methodology and Results

1.2.1 Part (a): Prior Probabilities

The prior probabilities, $P(Y = \text{cheetah})$ and $P(Y = \text{grass})$, were re-evaluated using the new `TrainingSamplesDCT_8_new.mat` data. The Maximum Likelihood Estimate (MLE) for the parameter p_c of a categorical distribution is given by:

$$\hat{p}_c = P(Y = c) = \frac{N_c}{N_{\text{total}}}$$

where N_c is the number of samples for class c . Based on the 250 foreground (cheetah) and 1053 background (grass) samples:

- $P(Y = \text{cheetah}) = \frac{250}{250+1053} = \frac{250}{1303} \approx 0.1919$
- $P(Y = \text{grass}) = \frac{1053}{250+1053} = \frac{1053}{1303} \approx 0.8081$

These results are identical to the priors computed in Quiz 1.

1.2.2 Part (b): Class-Conditional Parameters and Feature Selection

The class-conditional densities were modeled as 64-dimensional Gaussian distributions, $P_{X|Y}(x|c) \sim \mathcal{N}(\boldsymbol{\mu}_c, \boldsymbol{\Sigma}_c)$. The MLE parameters for the mean vector $\boldsymbol{\mu}_c$ and covariance matrix $\boldsymbol{\Sigma}_c$ were computed from the training data. Notably, the MLE for the covariance matrix uses a denominator of N (i.e., `ddof=0` in `numpy.cov`).

Instead of manual visual inspection, feature selection was performed quantitatively using the Symmetric Kullback-Leibler (KL) Divergence. This metric measures the "distance" between the 1D marginal distributions $P(X_k|\text{cheetah})$ and $P(X_k|\text{grass})$ for each of the 64 features (DCT coefficients). The 8 features with the highest KL divergence and the 8 with the lowest were identified.

- Best 8 Features (by KL Div): [0, 31, 26, 24, 39, 32, 17, 40]
- Worst 8 Features (by KL Div): [63, 62, 58, 2, 4, 61, 3, 59]

The marginal densities for these selected features are plotted in Figure 1.




Figure 1: Marginal densities $P(X_k|Y)$ for the 8 best (top) and 8 worst (bottom) features, as selected by Symmetric KL Divergence.

1.2.3 Part (c): Image Segmentation and Performance

The `cheetah.bmp` image was classified using a sliding 8x8 window. At each pixel, the 64-D DCT vector for the corresponding block was extracted and classified using the Bayes Decision Rule for minimum error:

$$\hat{Y} = \arg \max_{Y \in \{\text{cheetah, grass}\}} [\log P_{X|Y}(x|Y) + \log P(Y)]$$

This classification was performed twice:

1. Using the full 64-dimensional Gaussian models.
2. Using 8-dimensional Gaussian models built from the best 8 features.

The resulting probability of error for each classifier, computed by comparing the output mask to the ground truth, was:

- 64-D Classifier Error: 0.1450 (14.50%)
- 8-D Classifier Error: 0.0748 (7.48%)

1.2.4 Discussion: Explaining the Results

As shown by the error rates and the segmentation masks (Figures 2 and 3), the 8-dimensional classifier performs significantly better than the 64-dimensional one.

This result is a classic example of the Curse of Dimensionality. With a limited number of training samples (250 for cheetah, 1053 for grass), accurately estimating the parameters for a 64-dimensional Gaussian distribution is extremely difficult. The 64×64 covariance matrix has $\frac{64 \times 65}{2} = 2080$ unique

64D Gaussian (KL) (Error: 0.1450)




Figure 2: Segmentation using the full 64-D Gaussian model. Left: Predicted Mask. Right: Ground Truth. (Error: 14.50%)

8D Gaussian (KL) (Error: 0.0748)




Figure 3: Segmentation using the 8-D Gaussian model (best features). Left: Predicted Mask. Right: Ground Truth. This model achieves a lower error rate. (Error: 7.48%)

parameters. Estimating this many parameters from only 250 foreground samples leads to a model that is highly overfit to the training data and does not generalize well.

The 8-dimensional model, requires estimating an 8x8 covariance matrix (36 parameters). This model is simpler, its parameters can be estimated more robustly from the available data, and it consequently generalizes better to the unseen test image, resulting in a lower probability of error.

1.3 Appendix: Source Code

The Python source code used for this assignment is attached below.

```
1 import numpy as np
2 import scipy.io
3 import scipy.stats
4 import imageio.v3 as imageio
5 import matplotlib
6 import os
7 import warnings
8 from scipy.fftpack import dctn
9 from tqdm import tqdm
10
11 try:
12     matplotlib.use('Agg')
13 except ImportError:
14     pass
15
16 import matplotlib.pyplot as plt
17
18 warnings.filterwarnings('ignore')
19 np.random.seed(42)
20
21 BLOCK_SIZE = 8
22 DATA_FILE = 'data/TrainingSamplesDCT_8_new.mat'
23 ZIG_ZAG_FILE = 'data/Zig-Zag Pattern.txt'
24 IMAGE_FILE = 'data/cheetah.bmp'
25 MASK_FILE = 'data/cheetah_mask.bmp'
26 OUTPUT_DIR = 'hw2/output/'
27
28 def load_zig_zag_map(filepath):
29     """Loads the zig-zag scan pattern."""
30     zig_zag_pattern = np.loadtxt(filepath, dtype=int)
31     return np.argsort(zig_zag_pattern.flatten())
32
33
34 def dct2(block):
35     """Compute 2D DCT of an 8x8 block."""
36     return dctn(block, type=2, norm='ortho')
37
38
39 def compute_mle_parameters(data):
40     """
41     Compute MLE for mean and covariance matrix (ddof=0).
42     """
43     mean = np.mean(data, axis=0)
44     cov = np.cov(data, rowvar=False, ddof=0)
45     reg = np.eye(cov.shape[0]) * 1e-6
46     cov_reg = cov + reg
47     return mean, cov_reg
48
49
50 def compute_kl_divergence(p, q):
51     """Compute symmetric KL divergence between two 1D Gaussian distributions."""
52     mu_p, var_p = p
53     mu_q, var_q = q
54     epsilon = 1e-10
55     var_p = max(var_p, epsilon)
```

```

56     var_q = max(var_q, epsilon)
57
58     kl_pq = 0.5 * (np.log(var_q / var_p) + (var_p + (mu_p - mu_q) ** 2) / var_q - 1)
59     kl_qp = 0.5 * (np.log(var_p / var_q) + (var_q + (mu_q - mu_p) ** 2) / var_p - 1)
60
61     return kl_pq + kl_qp
62
63 def select_features_kl(fg_data, bg_data):
64     """Select best and worst 8 features based on symmetric KL divergence."""
65     kl_divergences = []
66     for i in range(64):
67         mean_fg = np.mean(fg_data[:, i])
68         var_fg = np.var(fg_data[:, i], ddof=0)
69         mean_bg = np.mean(bg_data[:, i])
70         var_bg = np.var(bg_data[:, i], ddof=0)
71         kl = compute_kl_divergence((mean_fg, var_fg), (mean_bg, var_bg))
72         kl_divergences.append(kl)
73
74     sorted_indices = np.argsort(kl_divergences)
75     worst_8_indices = sorted_indices[:8]
76     best_8_indices = sorted_indices[-8:][::-1]
77     return best_8_indices, worst_8_indices
78
79
80 def plot_best_worst_features(fg_data, bg_data, best_indices, worst_indices,
81                             output_path):
82     """Plot marginals for best and worst 8 features."""
83     fig, axes = plt.subplots(2, 8, figsize=(20, 7))
84     fig.suptitle("Marginal Densities (Best and Worst 8 Features by KL Div.)",
85                  fontsize=16)
86
87     for i, idx in enumerate(best_indices):
88         ax = axes[0, i]
89         mean_fg = np.mean(fg_data[:, idx])
90         std_fg = np.std(fg_data[:, idx], ddof=0)
91         mean_bg = np.mean(bg_data[:, idx])
92         std_bg = np.std(bg_data[:, idx], ddof=0)
93
94         x_min = min(mean_fg - 4 * std_fg, mean_bg - 4 * std_bg)
95         x_max = max(mean_fg + 4 * std_fg, mean_bg + 4 * std_bg)
96         x = np.linspace(x_min, x_max, 200)
97
98         ax.plot(x, scipy.stats.norm.pdf(x, mean_fg, std_fg), 'b-', label='Cheetah')
99     if i == 0: ax.legend()
100
101    for i, idx in enumerate(worst_indices):
102        ax = axes[1, i]
103        mean_fg = np.mean(fg_data[:, idx])
104        std_fg = np.std(fg_data[:, idx], ddof=0)
105        mean_bg = np.mean(bg_data[:, idx])
106        std_bg = np.std(bg_data[:, idx], ddof=0)
107
108        x_min = min(mean_fg - 4 * std_fg, mean_bg - 4 * std_bg)
109        x_max = max(mean_fg + 4 * std_fg, mean_bg + 4 * std_bg)

```

```

110     x = np.linspace(x_min, x_max, 200)
111
112     ax.plot(x, scipy.stats.norm.pdf(x, mean_fg, std_fg), 'b-')
113     ax.plot(x, scipy.stats.norm.pdf(x, mean_bg, std_bg), 'r-')
114     ax.set_title(f'Worst Feature {idx}', fontsize=10)
115
116 plt.tight_layout(rect=[0, 0.03, 1, 0.95])
117 plt.savefig(output_path, dpi=150)
118 plt.close(fig)
119
120
121 def plot_segmentation_results(predicted_mask, true_mask, title, output_path):
122     """Plot segmentation results."""
123     fig, axes = plt.subplots(1, 2, figsize=(12, 6))
124     axes[0].imshow(predicted_mask, cmap='gray')
125     axes[0].set_title('Predicted Mask', fontsize=14)
126     axes[0].axis('off')
127     axes[1].imshow(true_mask, cmap='gray')
128     axes[1].set_title('Ground Truth Mask', fontsize=14)
129     axes[1].axis('off')
130     plt.suptitle(title, fontsize=16)
131     plt.tight_layout(rect=[0, 0.03, 1, 0.95])
132     plt.savefig(output_path, dpi=150)
133     plt.close(fig)
134
135
136 def extract_dct_vectors_sliding_window(image, zig_zag_map):
137     """
138         Processes an image using a sliding window.
139         Returns a (H*W, 64) array of DCT vectors.
140     """
141
142     img = np.float32(image)
143     h, w = img.shape
144
145     img_padded = np.pad(img, ((0, BLOCK_SIZE - 1), (0, BLOCK_SIZE - 1)), mode='reflect')
146     all_dct_vectors = np.zeros((h * w, BLOCK_SIZE * BLOCK_SIZE))
147
148     idx = 0
149     for i in tqdm(range(h), desc="Extracting DCT"):
150         for j in range(w):
151             block = img_padded[i:i + BLOCK_SIZE, j:j + BLOCK_SIZE]
152             block_dct = dct2(block)
153             all_dct_vectors[idx] = block_dct.flatten()[zig_zag_map]
154             idx += 1
155
156     return all_dct_vectors, h, w
157
158 def classify_blocks_vectorized(X_test, mean_fg, cov_fg, mean_bg, cov_bg,
159                               log_prior_fg, log_prior_bg):
160     """
161         Classifies a set of test vectors using Bayesian decision rule.
162     """
163
164     log_ll_fg = scipy.stats.multivariate_normal.logpdf(X_test, mean=mean_fg, cov=cov_fg, allow_singular=True)
165     log_ll_bg = scipy.stats.multivariate_normal.logpdf(X_test, mean=mean_bg, cov=cov_bg, allow_singular=True)

```

```

165     g_fg = log_ll_fg + log_prior_fg
166     g_bg = log_ll_bg + log_prior_grass
167     decisions = (g_fg > g_bg).astype(int)
168     return decisions
169
170
171 def compute_error_rate(predicted_mask, true_mask, prior_fg, prior_bg):
172     """
173     Compute Bayesian probability of error (weighted).
174     P(error) = P(error|cheetah)P(cheetah) + P(error|grass)P(grass)
175     """
176     predicted_binary = (predicted_mask > 0.5).astype(int)
177     true_binary = (true_mask > 0.5).astype(int)
178
179     fg_pixels_total = np.sum(true_binary == 1)
180     bg_pixels_total = np.sum(true_binary == 0)
181
182     p_error_given_fg = 0.0
183     if fg_pixels_total > 0:
184         fg_misclassified = np.sum((true_binary == 1) & (predicted_binary == 0))
185         p_error_given_fg = fg_misclassified / fg_pixels_total
186
187     p_error_given_bg = 0.0
188     if bg_pixels_total > 0:
189         bg_misclassified = np.sum((true_binary == 0) & (predicted_binary == 1))
190         p_error_given_bg = bg_misclassified / bg_pixels_total
191
192     total_error = (p_error_given_fg * prior_fg) + (p_error_given_bg * prior_bg)
193     return total_error
194
195
196 # --- Main Execution ---
197
198 def main():
199     os.makedirs(OUTPUT_DIR, exist_ok=True)
200
201     # 1. Load Data
202     zig_zag_map = load_zig_zag_map(ZIG_ZAG_FILE)
203     train_data = scipy.io.loadmat(DATA_FILE)
204     fg_data = train_data['TrainsampleDCT_FG']
205     bg_data = train_data['TrainsampleDCT_BG']
206
207     image = imageio.imread(IMAGE_FILE, mode='L')
208     true_mask = imageio.imread(MASK_FILE)
209     true_mask = (true_mask > 127).astype(int)
210
211
212     # 2a. Priors
213     n_fg = fg_data.shape[0]
214     n_bg = bg_data.shape[0]
215     n_total = n_fg + n_bg
216     prior_cheetah = n_fg / n_total
217     prior_grass = n_bg / n_total
218     log_prior_cheetah = np.log(prior_cheetah)
219     log_prior_grass = np.log(prior_grass)
220
221     print("\n" + "=" * 70)
222     print("Problem 6(a): Prior Probabilities")
223     print(f"P(Y=cheetah): {prior_cheetah:.6f} (N={n_fg})")

```

```

224 print(f"P(Y=grass): {prior_grass:.6f} (N={n_bg})")
225
226 # 2b. ML Parameters & Feature Selection
227 print("\n" + "=" * 70)
228 print("Problem 6(b): ML Parameters and Feature Selection")
229
230 mean_fg_64, cov_fg_64 = compute_mle_parameters(fg_data)
231 mean_bg_64, cov_bg_64 = compute_mle_parameters(bg_data)
232
233 best_8_indices, worst_8_indices = select_features_kl(fg_data, bg_data)
234
235 print(f"Best 8 features (KL Div): {best_8_indices.tolist()}")
236 print(f"Worst 8 features (KL Div): {worst_8_indices.tolist()}")
237
238 plot_features_filename = os.path.join(OUTPUT_DIR, 'best_worst_8_features_KL.
png')
239 plot_best_worst_features(fg_data, bg_data, best_8_indices, worst_8_indices,
240                         plot_features_filename)
241 print(f"Feature plots saved to {plot_features_filename}")
242
243
244 # 3. Process Test Image
245 print("\n" + "=" * 70)
246 print("Problem 6(c): Classification (Sliding Window)")
247 X_test_64, img_h, img_w = extract_dct_vectors_sliding_window(image,
248 zig_zag_map)
249
250
251 # 4. Classify 64D
252 print("Classifying with 64D Gaussians...")
253 decisions_64_flat = classify_blocks_vectorized(X_test_64,
254                                                 mean_fg_64, cov_fg_64,
255                                                 mean_bg_64, cov_bg_64,
256                                                 log_prior_cheetah,
257                                                 log_prior_grass)
258 mask_64d = decisions_64_flat.reshape(img_h, img_w)
259 error_64d = compute_error_rate(mask_64d, true_mask, prior_cheetah, prior_grass
)
260
261 plot_64d_filename = os.path.join(OUTPUT_DIR, 'segmentation_64d_KL.png')
262 plot_segmentation_results(mask_64d, true_mask,
263                            f'64D Gaussian (KL) (Error: {error_64d:.4f})',
264                            plot_64d_filename)
265 print(f"64D plot saved to {plot_64d_filename}")
266
267 # 5. Classify 8D
268 print("Classifying with 8D Gaussians (best features)...")
269 X_test_8 = X_test_64[:, best_8_indices]
270
271 mean_fg_8 = mean_fg_64[best_8_indices]
272 mean_bg_8 = mean_bg_64[best_8_indices]
273 cov_fg_8 = cov_fg_64[np.ix_(best_8_indices, best_8_indices)]
274 cov_bg_8 = cov_bg_64[np.ix_(best_8_indices, best_8_indices)]
275
276 decisions_8_flat = classify_blocks_vectorized(X_test_8,
277                                                 mean_fg_8, cov_fg_8,
278                                                 mean_bg_8, cov_bg_8,
279                                                 log_prior_cheetah,
280                                                 log_prior_grass)

```

```

278 mask_8d = decisions_8_flat.reshape(img_h, img_w)
279 error_8d = compute_error_rate(mask_8d, true_mask, prior_cheetah, prior_grass)
280
281 plot_8d_filename = os.path.join(OUTPUT_DIR, 'segmentation_8d_KL.png')
282 plot_segmentation_results(mask_8d, true_mask,
283                             f'8D Gaussian (KL) (Error: {error_8d:.4f})',
284                             plot_8d_filename)
285 print(f"8D plot saved to {plot_8d_filename}")
286
287 # 6. Final Results
288 print("\n" + "=" * 70)
289 print("Final Explanation of Results")
290 print("=" * 70)
291 print(f"64D Classifier Bayesian Error: {error_64d:.6f}")
292 print(f" 8D Classifier Bayesian Error: {error_8d:.6f}")
293
294 if error_8d < error_64d:
295     print("\nSUCCESS: The 8-dimensional classifier performed BETTER (lower
296          error).")
297     print("This confirms the 'Curse of Dimensionality'.")
298     print(f"With limited samples (N_fg={n_fg}, N_bg={n_bg}), estimating a 64
299          x64")
300     print("covariance matrix is unstable. The 8D model is simpler and
301          generalizes better.")
302 else:
303     print("\nNOTE: The 64D classifier performed better or equal to the 8D
304          classifier.")
305
306 if __name__ == "__main__":
307     main()

```

Listing 1: Python code for HW2 (hw2/hw2_solution.py)