

ECE 271A: Statistical Learning I Quiz Report

JunPyung Kim
PID: A69044427
University of California, San Diego

December 17, 2025

1 Quiz 3: Bayesian Parameter Estimation

1.1 Objective

The objective of this assignment is to explore Bayesian Parameter Estimation by treating the parameters of the class-conditional densities as random variables. We compare three different estimators:

1. **Maximum Likelihood (ML) Solution:** Computes parameters solely from the training data.
2. **Maximum A Posteriori (MAP) Solution:** Estimates the parameter vector that maximizes the posterior density.
3. **Bayesian Predictive Solution:** Marginals out the unknown parameters to compute the predictive distribution $P(x|\mathcal{D})$.

We analyze the performance of these estimators as a function of the prior uncertainty (α) and the training set size (N), using two different strategies for the prior mean.

1.2 Methodology

1.2.1 Model Setup

We model the class-conditional densities as multivariate Gaussians. For this assignment, we assume the covariance Σ is known (approximated by the sample covariance of the training set), but the mean μ is unknown. The prior distribution for the mean is modeled as a Gaussian:

$$P(\mu) = \mathcal{G}(\mu, \mu_0, \Sigma_0)$$

where $\Sigma_0 = \alpha \mathbf{W}$ and \mathbf{W} is diagonal (i.e., $\Sigma_0 = \alpha \text{diag}(\mathbf{W}_0)$). The parameter α scales the variance of the prior; a small α indicates high confidence in the prior mean μ_0 , while a large α indicates high uncertainty.

Decision rule and evaluation. Class priors in the decision rule are the ML estimates from the training set. The Probability of Error (PoE) is computed on the *cheetah test image* as the *pixel-wise misclassification rate* using the provided ground-truth mask.

1.2.2 Strategies

Two strategies were used to select the prior mean μ_0 (provided in `Prior_1.mat` and `Prior_2.mat`):

- **Strategy 1 (More informative prior):** μ_0 is closer to the training-data ML means for each class.
- **Strategy 2 (Less informative / poorer prior):** μ_0 is farther from the ML means (more generic).

1.3 Results

The Probability of Error (PoE) was computed for varying values of α across four datasets (\mathcal{D}_1 to \mathcal{D}_4) of increasing size. In all figures, the ML curve is constant w.r.t. α , and the MAP/Predictive curves converge to the ML baseline as α becomes large.

1.3.1 Strategy 1: More informative prior

Figures 1 and 2 show the results for Strategy 1. For small α , MAP and Bayesian Predictive solutions are influenced strongly by the prior. Depending on how well μ_0 aligns with the test-image distribution, this can slightly improve or slightly degrade PoE relative to ML. As α increases, the influence of the prior diminishes and both MAP and Predictive approaches approach the ML baseline.




Figure 1: Strategy 1 Results for Dataset 1 (Left) and Dataset 2 (Right), using $d = 64$ DCT coefficients.




Figure 2: Strategy 1 Results for Dataset 3 (Left) and Dataset 4 (Right), using $d = 64$ DCT coefficients.

1.3.2 Strategy 2: Less informative / poorer prior

Figures 3 and 4 show the results for Strategy 2. With a less accurate prior mean, very small α can pull the posterior mean away from the data-driven estimate and increase PoE. However, for small training sets, a strongly regularizing prior can sometimes reduce variance and yield comparable (or occasionally slightly better) performance than ML at certain α values. As α increases, the curves again converge to the ML baseline.




Figure 3: Strategy 2 Results for Dataset 1 (Left) and Dataset 2 (Right), using $d = 64$ DCT coefficients.

1.4 Discussion

1.4.1 Relative Behavior of Estimators

The ML solution appears as a flat line because it does not depend on the prior parameter α . The MAP and Bayesian Predictive solutions vary with α :

- When $\alpha \rightarrow 0$, the prior dominates and the posterior mean is pulled toward μ_0 . The resulting PoE can be lower or higher than ML depending on how well the prior matches the true/test distribution.




Figure 4: Strategy 2 Results for Dataset 3 (Left) and Dataset 4 (Right), using $d = 64$ DCT coefficients.

- When $\alpha \rightarrow \infty$, the prior becomes effectively uninformative, and the MAP estimate of the mean approaches the ML mean (thus the MAP curve approaches the ML baseline).

The Bayesian Predictive classifier typically tracks the MAP curve closely, but can differ slightly because it incorporates posterior uncertainty in μ (through an added covariance term in the predictive distribution).

1.4.2 Effect of Dataset Size

Comparing \mathcal{D}_1 (smallest) to \mathcal{D}_4 (largest):

- **Higher variance in small datasets:** With fewer training samples, the ML mean estimates are noisier, which generally increases PoE (e.g., the ML baseline is higher on smaller datasets).
- **Prior influence lasts longer for small datasets:** The posterior mean balances μ_0 vs. the sample mean with a term proportional to Σ/N ; smaller N makes the data term less concentrated, so the prior can influence the solution over a wider range of α .
- **Faster convergence for larger datasets:** With larger N , the curves converge to ML at smaller α values since the data dominates the posterior more quickly.

1.5 Appendix: Source Code

The Python source code used to generate these results is included below (`hw3/hw3_solution.py`).

```
1 #!/usr/bin/env python3
2 import os
3 import warnings
4 from pathlib import Path
5
6 import numpy as np
7 import scipy.io
8 import imageio.v3 as imageio
9 import matplotlib
10 from scipy.fftpack import dctn
11 from tqdm import tqdm
12
13 try:
14     matplotlib.use("Agg")
15 except Exception:
16     pass
17
18 import matplotlib.pyplot as plt
19 warnings.filterwarnings("ignore")
20
21 BLOCK_SIZE = 8
22
23 THIS_DIR = Path(__file__).resolve().parent          # .../hw3
24 PROJECT_ROOT = THIS_DIR.parent                     # .../ece271a
25 DATA_DIR = PROJECT_ROOT / "data"
26 OUTPUT_DIR = THIS_DIR / "output"
27
28 CHEETAH_IMG = DATA_DIR / "cheetah.bmp"
29 CHEETAH_MASK = DATA_DIR / "cheetah_mask.bmp"
30 ZIG_ZAG_FILE = DATA_DIR / "Zig-Zag Pattern.txt"
31 OUTPUT_DIR.mkdir(parents=True, exist_ok=True)
32
33 def load_mat_file(filename: str) -> dict:
34     return scipy.io.loadmat(str(DATA_DIR / filename))
35
36 def load_zig_zag_map(filepath: Path) -> np.ndarray:
37     zig_zag_pattern = np.loadtxt(str(filepath), dtype=int)
38     return np.argsort(zig_zag_pattern.flatten())
39
40 def dct2(block: np.ndarray) -> np.ndarray:
41     return dctn(block, type=2, norm="ortho")
42
43 def extract_dct_features(image_path: Path, zig_zag_map: np.ndarray):
44     img = imageio.imread(str(image_path), mode="L").astype(np.float32) / 255.0
45     h, w = img.shape
46     img_padded = np.pad(img, ((0, BLOCK_SIZE - 1), (0, BLOCK_SIZE - 1)), mode="constant")
47
48     num_vectors = h * w
49     features = np.zeros((num_vectors, 64), dtype=np.float32)
50
51     idx = 0
52     for r in range(h):
53         for c in range(w):
54             block = img_padded[r : r + BLOCK_SIZE, c : c + BLOCK_SIZE]
```

```

55     dct_block = dct2(block).flatten()
56     features[idx] = dct_block[zig_zag_map]
57     idx += 1
58   return features, h, w
59
60 def binarize_mask(mask_flat: np.ndarray) -> np.ndarray:
61   return (mask_flat >= 128).astype(np.int32)
62
63 def compute_poe_pixelwise(pred_labels: np.ndarray, gt_labels01: np.ndarray) ->
64   float:
65   return float(np.mean(pred_labels != gt_labels01))
66
67 def gaussian_log_likelihood_batch(X: np.ndarray, mu: np.ndarray, cov: np.ndarray,
68   eps: float = 1e-8) -> np.ndarray:
69   d = X.shape[1]
70   cov = cov + np.eye(d) * eps
71   sign, logdet = np.linalg.slogdet(cov)
72   if sign <= 0:
73     logdet = np.log(np.linalg.det(cov) + 1e-30)
74   inv_cov = np.linalg.inv(cov)
75
76   diff = X - mu
77   mahal = np.sum((diff @ inv_cov) * diff, axis=1)
78   const = -0.5 * (d * np.log(2 * np.pi) + logdet)
79   return const - 0.5 * mahal
80
81 def posterior_for_mean(mu0: np.ndarray, Sigma0: np.ndarray, X: np.ndarray):
82   N = X.shape[0]
83   xbar = np.mean(X, axis=0)
84   Sigma = np.cov(X, rowvar=False, ddof=0)
85
86   d = X.shape[1]
87   Sigma = Sigma + np.eye(d) * 1e-8
88   Sigma_over_N = Sigma / max(N, 1)
89
90   A = Sigma0 + Sigma_over_N
91   A_inv = np.linalg.inv(A)
92
93   Sigma_n = Sigma0 @ A_inv @ Sigma_over_N
94   mu_n = (Sigma0 @ A_inv @ xbar) + (Sigma_over_N @ A_inv @ mu0)
95   return mu_n, Sigma_n
96
97 def solve():
98   print("[INFO] Loading data...")
99   zig_zag = load_zig_zag_map(ZIG_ZAG_FILE)
100
101   alpha_vec = load_mat_file("Alpha.mat")["alpha"].flatten()
102   train_subsets = load_mat_file("TrainingSamplesDCT_subsets_8.mat")
103   prior_1 = load_mat_file("Prior_1.mat")
104   prior_2 = load_mat_file("Prior_2.mat")
105
106   print("[INFO] Processing test image / mask...")
107   test_features_64, _, _ = extract_dct_features(CHEETAH_IMG, zig_zag)
108   gt_mask_flat = imageio.imread(str(CHEETAH_MASK)).flatten()
109   gt01 = binarize_mask(gt_mask_flat)
110
111   strategies = [
112     {"name": "Strategy 1", "data": prior_1},
113     {"name": "Strategy 2", "data": prior_2},

```

```

112     ]
113
114     datasets = [
115         {"id": 1, "bg": train_subsets["D1_BG"], "fg": train_subsets["D1_FG"]},
116         {"id": 2, "bg": train_subsets["D2_BG"], "fg": train_subsets["D2_FG"]},
117         {"id": 3, "bg": train_subsets["D3_BG"], "fg": train_subsets["D3_FG"]},
118         {"id": 4, "bg": train_subsets["D4_BG"], "fg": train_subsets["D4_FG"]},
119     ]
120
121     dims_to_run = [64]
122
123     print("[INFO] Starting classification loop...")
124     for strat in strategies:
125         strat_name = strat["name"]
126         strat_data = strat["data"]
127
128         W0_64 = strat_data["W0"].flatten()
129         mu0_FG_64 = strat_data["mu0_FG"].flatten()
130         mu0_BG_64 = strat_data["mu0_BG"].flatten()
131
132         print(f"\n--- {strat_name} ---")
133
134         for d in dims_to_run:
135             test_features = test_features_64[:, :d]
136             W0 = W0_64[:d]
137             mu0_FG = mu0_FG_64[:d]
138             mu0_BG = mu0_BG_64[:d]
139
140             for dataset in datasets:
141                 d_id = dataset["id"]
142                 fg_train = dataset["fg"][:, :d]
143                 bg_train = dataset["bg"][:, :d]
144
145                 n_fg = fg_train.shape[0]
146                 n_bg = bg_train.shape[0]
147
148                 p_fg = n_fg / (n_fg + n_bg)
149                 p_bg = n_bg / (n_fg + n_bg)
150                 log_p_fg = np.log(p_fg + 1e-30)
151                 log_p_bg = np.log(p_bg + 1e-30)
152
153                 mu_ml_fg = np.mean(fg_train, axis=0)
154                 cov_ml_fg = np.cov(fg_train, rowvar=False, ddof=0) + np.eye(d) * 1
155
156                 mu_ml_bg = np.mean(bg_train, axis=0)
157                 cov_ml_bg = np.cov(bg_train, rowvar=False, ddof=0) + np.eye(d) * 1
158
159                 ll_fg_ml = gaussian_log_likelihood_batch(test_features, mu_ml_fg,
160                 cov_ml_fg)
161                 ll_bg_ml = gaussian_log_likelihood_batch(test_features, mu_ml_bg,
162                 cov_ml_bg)
163                 pred_ml = ((ll_fg_ml + log_p_fg) > (ll_bg_ml + log_p_bg)).astype(
164                 np.int32)
165
166                 poe_ml = compute_poe_pixelwise(pred_ml, gt01)
167
168                 errors_ml, errors_map, errors_bayes = [], [], []

```

```

165         print(f"[INFO] {strat_name} | Dataset {d_id} | d={d}: sweeping
166         alpha...")
167         for alpha in tqdm(alpha_vec, leave=False):
168             Sigma0 = np.diag(alpha * W0 + 1e-30)
169
170             mu_n_fg, Sigma_n_fg = posterior_for_mean(mu0_FG, Sigma0,
171             fg_train)
172             mu_n_bg, Sigma_n_bg = posterior_for_mean(mu0_BG, Sigma0,
173             bg_train)
174
175             ll_fg_map = gaussian_log_likelihood_batch(test_features,
176             mu_n_fg, cov_ml_fg)
177             ll_bg_map = gaussian_log_likelihood_batch(test_features,
178             mu_n_bg, cov_ml_bg)
179             pred_map = ((ll_fg_map + log_p_fg) > (ll_bg_map + log_p_bg)).
180             astype(np.int32)
181             poe_map = compute_poe_pixelwise(pred_map, gt01)
182
183             pred_cov_fg = cov_ml_fg + Sigma_n_fg
184             pred_cov_bg = cov_ml_bg + Sigma_n_bg
185             ll_fg_bayes = gaussian_log_likelihood_batch(test_features,
186             mu_n_fg, pred_cov_fg)
187             ll_bg_bayes = gaussian_log_likelihood_batch(test_features,
188             mu_n_bg, pred_cov_bg)
189             pred_bayes = ((ll_fg_bayes + log_p_fg) > (ll_bg_bayes +
190             log_p_bg)).astype(np.int32)
191             poe_bayes = compute_poe_pixelwise(pred_bayes, gt01)
192
193             errors_ml.append(poe_ml)
194             errors_map.append(poe_map)
195             errors_bayes.append(poe_bayes)
196
197             plt.figure(figsize=(10, 6))
198             plt.semilogx(alpha_vec, errors_bayes, label="Predictive (Bayesian)",
199             ", linewidth=2")
200             plt.semilogx(alpha_vec, errors_map, label="MAP", linewidth=2,
201             linestyle="--")
202             plt.semilogx(alpha_vec, errors_ml, label="ML", linewidth=2,
203             linestyle="-.")
204
205             plt.title(f"PoE vs Alpha - {strat_name} - Dataset {d_id} (d={d})")
206             plt.xlabel("Alpha")
207             plt.ylabel("Probability of Error (pixel-wise on cheetah image)")
208             plt.legend()
209             plt.grid(True, which="both", ls="--", alpha=0.4)
210
211             plot_filename = f"PoE_Strategy_{strat_name[-1]}_Dataset_{d_id}_d{d}.
212             png"
213             plt.savefig(str(OUTPUT_DIR / plot_filename), dpi=150, bbox_inches=
214             "tight")
215             plt.close()
216
217             print(f"\n[DONE] Plots saved to: {OUTPUT_DIR}")
218
219 if __name__ == "__main__":
220     solve()

```

Listing 1: Python code for HW3 (final implementation)