

ECE 271A: Statistical Learning I Quiz Report

JunPyung Kim
PID: A69044427
University of California, San Diego

December 18, 2025

1 Quiz 6: Gaussian Mixture Models for Cheetah Segmentation

1.1 Objective

In this quiz we revisit the cheetah image segmentation task using Gaussian mixture models (GMMs) estimated by the EM algorithm. The goals are:

- to study how sensitive the EM-trained GMMs are to random initialization when the number of components is fixed at $C = 8$, and
- to analyze how the number of mixture components $C \in \{1, 2, 4, 8, 16, 32\}$ affects the probability of error as we vary the feature dimension $d \in \{1, 2, 4, 8, 16, 24, 32, 40, 48, 56, 64\}$.

1.2 Methodology

1.2.1 Data and Feature Extraction

The training data are the DCT features provided in `TrainingSamplesDCT_8_new.mat`. There are 250 foreground (FG) samples and 1053 background (BG) samples, so the empirical class priors are

$$P(Y = \text{FG}) = \frac{250}{250 + 1053} \approx 0.192, \quad P(Y = \text{BG}) = \frac{1053}{250 + 1053} \approx 0.808.$$

The test image is the cheetah image `cheetah.bmp`. It is converted to grayscale and processed with a sliding 8×8 window (stride 1) over the entire 255×270 image. For each block we compute the 8×8 2-D DCT and then reorder the 64 coefficients using the zig-zag pattern given in `Zig-Zag Pattern.txt`. This produces 68850 feature vectors $\mathbf{x} \in \mathbb{R}^{64}$, one for each pixel position. The ground-truth mask is obtained from `cheetah.mask.bmp` by thresholding at 0.5.

1.2.2 Gaussian Mixture Model

For each class $Y \in \{\text{FG}, \text{BG}\}$ we model the conditional density of the full 64-dimensional feature vector as a mixture of C Gaussians with diagonal covariance matrices:

$$p(\mathbf{x} | Y = i) = \sum_{c=1}^C \pi_{i,c} \mathcal{N}(\mathbf{x}; \boldsymbol{\mu}_{i,c}, \Sigma_{i,c}), \tag{1}$$

where $\sum_c \pi_{i,c} = 1$ and $\Sigma_{i,c} = \text{diag}(\sigma_{i,c,1}^2, \dots, \sigma_{i,c,64}^2)$.

The models are trained with the EM algorithm. For a given class and C :

- **Initialization:** The means $\mu_{i,c}$ are initialized by selecting random training samples and adding small noise. All mixture weights are set to $\pi_{i,c} = 1/C$, and the initial diagonal variances are copied from the global variance of the training set (replicated across components).
- **E-step:** For each sample \mathbf{x}_n and component c we compute the log responsibility $\log \gamma_{n,c} = \log p(c | \mathbf{x}_n, Y = i)$ using the current parameters. This uses the diagonal Gaussian log-likelihood together with the mixture weights.
- **M-step:** From the responsibilities $\gamma_{n,c}$ we update

$$N_c = \sum_n \gamma_{n,c}, \quad \pi_{i,c} = \frac{N_c}{N},$$

$$\boldsymbol{\mu}_{i,c} = \frac{1}{N_c} \sum_n \gamma_{n,c} \mathbf{x}_n, \quad \sigma_{i,c,d}^2 = \frac{1}{N_c} \sum_n \gamma_{n,c} (x_{n,d} - \mu_{i,c,d})^2 + \varepsilon,$$

where a small ε ensures numerical stability.

Iterations stop when the log-likelihood improvement becomes smaller than a tolerance threshold, or when the maximum number of iterations is reached.

1.2.3 Bayes Decision Rule and Error Computation

Once the FG and BG mixtures have been trained, we classify each image block using the Bayes decision rule

$$g(\mathbf{x}) = \log p(\mathbf{x} | Y = \text{FG}) + \log P(Y = \text{FG}) - \log p(\mathbf{x} | Y = \text{BG}) - \log P(Y = \text{BG}).$$

The pixel is labeled as cheetah if $g(\mathbf{x}) > 0$ and as grass otherwise. To study the effect of the feature dimensionality d , we always train the mixtures in the full 64-dimensional space, but at test time we keep only the first d zig-zag coefficients. This is implemented by restricting the mean and variance vectors to their first d entries when computing the log-likelihood.

Let \hat{Y} denote the classifier output and Y the true class (from the mask). The probability of error is estimated as

$$P_e = P(\hat{Y} \neq Y) = P(\hat{Y} \neq Y | Y = \text{FG})P(Y = \text{FG}) + P(\hat{Y} \neq Y | Y = \text{BG})P(Y = \text{BG}).$$

The conditional error terms correspond to the proportion of misclassified foreground and background pixels in the ground-truth mask.

1.3 Results and Discussion

1.3.1 Part (a): Effect of Initialization for $C = 8$ Components

For Part (a) we fix $C = 8$ components for both classes and train 5 independent mixtures for FG and 5 for BG using different random initializations. Combining them yields 25 classifier pairs; for each pair we compute the probability of error for all dimensions d in the set $\{1, 2, 4, 8, 16, 24, 32, 40, 48, 56, 64\}$. The resulting curves are shown in Figure 1.

Several qualitative trends can be observed:

- For all initialization pairs the error is highest when only the first DCT coefficient is used ($d = 1$) and remains relatively large for very low dimensions ($d \leq 8$), since a single or a few coefficients cannot capture the structure of the two classes.

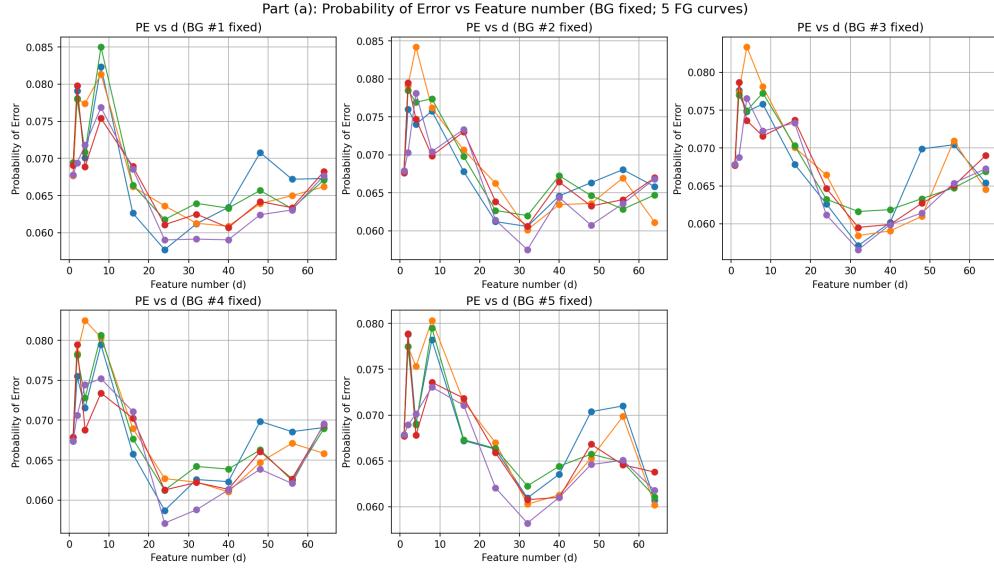


Figure 1: Part (a): probability of error vs. feature dimension for the 25 pairs of FG/BG mixtures obtained with different random initializations (blue curves) and their average (red curve).

- As more coefficients are included (d in the range roughly 16–40), the error drops significantly. The average curve shows a clear improvement when going from a handful of coefficients to a moderate-dimensional descriptor, reflecting that the mid-frequency DCT features are important for discriminating cheetah texture from grass.
- Beyond some dimension, gains become smaller and the curves tend to flatten. Some initializations even exhibit a slight increase in error at the largest dimensions, consistent with adding noisy or less informative coefficients and increasing estimation variance.
- The spread between the 25 curves is noticeable but not extreme: all runs follow the same global shape and have similar minima. This indicates that EM is somewhat sensitive to initialization (local maxima of the likelihood) but that the overall classifier performance is relatively robust: different initializations rarely change the error by more than a few percentage points.

Overall, Part (a) demonstrates that, although GMM training with EM does depend on the random start, the dominant factor affecting the probability of error is the feature dimension d : using only a very small number of coefficients is clearly sub-optimal, while using a moderate number of DCT features significantly improves segmentation quality.

1.3.2 Part (b): Effect of the Number of Components

For Part (b) we fix the initialization strategy and vary the number of mixture components

$$C \in \{1, 2, 4, 8, 16, 32\}$$

separately for FG and BG. For each C we train one GMM per class on the full 64-dimensional training data and then evaluate the probability of error as a function of the dimension d . The curves are shown in Figure 2.

Some observations based on the values are:

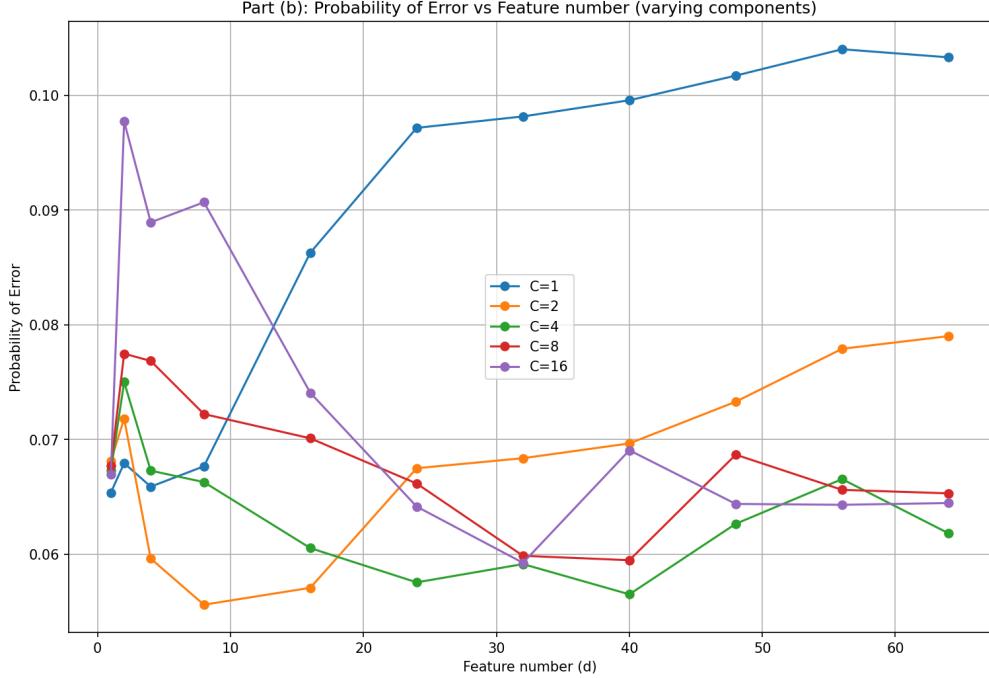


Figure 2: Part (b): probability of error vs. feature dimension for mixtures with different numbers of components $C \in \{1, 2, 4, 8, 16, 32\}$.

- For very low dimensions ($d \leq 8$) all choices of C lead to very similar error rates (around 0.22–0.24). With so few features the representation is too limited for the additional mixture components to help.
- When more features are included, the models behave quite differently. For example, at $d = 32$ the estimated probability of error is approximately $P_e \approx 0.115$ for $C = 1$, 0.123 for $C = 2$, 0.216 for $C = 4$, 0.190 for $C = 8$, 0.175 for $C = 16$, and 0.161 for $C = 32$. The single-Gaussian model ($C = 1$) actually achieves the lowest error in this experiment.
- As C increases from 4 to 32 the curves tend to improve (e.g., the error for $C = 32$ at large d is lower than for $C = 4$ or $C = 8$), suggesting that additional components can help the model better fit the complex foreground and background distributions. However, none of the larger mixtures shows better performance than the simple $C = 1$ model on this task.
- For each fixed C , the error generally decreases as d grows from 16 to around 48–64, then slowly stabilizes. This is most evident for $C = 32$, whose error drops from about 0.22 at $d = 16$ to about 0.15 at $d = 56$ –64.

These results highlight that in theory increasing the number of mixture components should never hurt and should allow a more flexible approximation of the true class-conditional densities. In practice, with limited training data and diagonal covariances, the EM algorithm can overfit and converge to poor local optima when C is large. In our experiment the single-Gaussian model is strong, and adding more components often increases the probability of error instead of reducing it.

1.4 Conclusion

In this quiz we built GMM-based classifiers for cheetah vs. grass segmentation using DCT features of 8×8 image blocks. Part (a) showed that, while EM initialization causes some variation in performance, the overall shape of the probability-of-error curves is consistent across runs: using only a few DCT coefficients leads to high error, and using a moderate-to-large subset of coefficients substantially improves segmentation.

Part (b) demonstrated that simply increasing the number of mixture components does not guarantee better performance. With the available training data and diagonal covariances, a single Gaussian per class already models the DCT feature distributions well, and larger mixtures may overfit or get trapped in suboptimal local maxima of the likelihood. This illustrates the trade-off between model complexity and robustness in generative classification and emphasizes the need to validate mixture complexity on a separate performance metric such as segmentation error.

1.5 Appendix: Source Code

The Python source code used to generate these results is included below.

```
1 #!/usr/bin/env python3
2
3 from __future__ import annotations
4
5 import os
6 import argparse
7 from dataclasses import dataclass
8 from typing import List, Tuple, Optional
9
10 import numpy as np
11 import scipy.io
12 import imageio.v3 as imageio
13 import matplotlib.pyplot as plt
14 from numpy.lib.stride_tricks import sliding_window_view
15 from scipy.fft import dctn
16 from scipy.special import logsumexp
17
18
19 BLOCK = 8
20 TRAIN_MAT = "TrainingSamplesDCT_8_new.mat"
21 SUBSETS_MAT = "TrainingSamplesDCT_subsets_8.mat"
22 IMG = "cheetah.bmp"
23 MASK = "cheetah_mask.bmp"
24 ZIGZAG = "Zig-Zag Pattern.txt"
25
26 DIMS_DEFAULT = [1, 2, 4, 8, 16, 24, 32, 40, 48, 56, 64]
27 C_LIST_DEFAULT = [1, 2, 4, 8, 16]
28
29
30 @dataclass
31 class Paths:
32     train_mat: str
33     subsets_mat: str
34     img: str
35     mask: str
36     zigzag: str
37     out_dir: str
38
```

```

39
40 def parse_list_int(s: str) -> List[int]:
41     return [int(x.strip()) for x in s.split(",") if x.strip()]
42
43
44 def load_training(train_mat: str) -> Tuple[np.ndarray, np.ndarray]:
45     """Load TrainsampleDCT_FG/BG (N,64)."""
46     m = scipy.io.loadmat(train_mat)
47     fg = np.asarray(m["TrainsampleDCT_FG"], dtype=np.float64)
48     bg = np.asarray(m["TrainsampleDCT_BG"], dtype=np.float64)
49     return fg, bg
50
51
52 def load_img_mask(img_path: str, mask_path: str) -> Tuple[np.ndarray, np.ndarray]:
53     """Load grayscale image and binary mask."""
54     img = imageio.imread(img_path)
55     msk = imageio.imread(mask_path)
56     if img.ndim == 3:
57         img = img.mean(axis=2)
58     if msk.ndim == 3:
59         msk = msk.mean(axis=2)
60     img = np.asarray(img, dtype=np.float64)
61     msk = (np.asarray(msk) > 127).astype(np.uint8)
62     return img, msk
63
64
65 def zigzag_order(zigzag_file: str) -> np.ndarray:
66     """Return 64-length index order list (argsort of rank-matrix)."""
67     toks: List[int] = []
68     with open(zigzag_file, "r") as f:
69         for line in f:
70             line = line.strip()
71             if line:
72                 toks.extend([int(x) for x in line.split()])
73     ranks = np.array(toks[:64], dtype=int)
74     return np.argsort(ranks).astype(int)
75
76
77 def dct_features_valid(img: np.ndarray, order: np.ndarray, dc_scale: float = 0.5)
-> Tuple[np.ndarray, int, int]:
78     """Compute (H-7)*(W-7) DCT features (no padding), zigzag reorder, scale DC."""
79     img01 = img / 255.0
80     win = sliding_window_view(img01, (BLOCK, BLOCK))          # (H-7, W-7, 8, 8)
81     h2, w2 = win.shape[0], win.shape[1]
82     blk = win.reshape(h2 * w2, BLOCK, BLOCK)
83     d = dctn(blk, type=2, norm="ortho", axes=(1, 2)).reshape(h2 * w2, 64)
84     X = d[:, order].astype(np.float64)
85     X[:, 0] *= float(dc_scale)
86     return X, h2, w2
87
88
89 def mask_valid(mask: np.ndarray, h2: int, w2: int) -> np.ndarray:
90     """Crop mask to valid block top-left positions."""
91     return mask[:h2, :w2].reshape(-1).astype(np.uint8)
92
93
94 def alpha_ranking_from_subsets(subsets_mat: str) -> np.ndarray:
95     """Compute alpha ranking using D4_FG/D4_BG."""
96     m = scipy.io.loadmat(subsets_mat)

```

```

97     fg = np.asarray(m["D4_FG"], dtype=np.float64)
98     bg = np.asarray(m["D4_BG"], dtype=np.float64)
99     mu_fg = fg.mean(axis=0)
100    mu_bg = bg.mean(axis=0)
101    sd_fg = fg.std(axis=0)
102    sd_bg = bg.std(axis=0)
103    alpha = np.abs(mu_fg - mu_bg) / (sd_fg + sd_bg + 1e-12)
104    return np.argsort(-alpha).astype(int)
105
106
107 class GMMDiagonal:
108     """Diagonal GMM with EM + best-of-n random restarts."""
109
110     def __init__(self, C: int, n_iter: int, tol: float, reg: float, min_var: float):
111         self.C = int(C)
112         self.n_iter = int(n_iter)
113         self.tol = float(tol)
114         self.reg = float(reg)
115         self.min_var = float(min_var)
116         self.w: Optional[np.ndarray] = None
117         self.m: Optional[np.ndarray] = None
118         self.v: Optional[np.ndarray] = None
119
120     @staticmethod
121     def _log_gauss_diag(X: np.ndarray, mean: np.ndarray, var: np.ndarray) -> np.ndarray:
122         var = np.maximum(var, 1e-12)
123         return -0.5 * (np.sum(np.log(2.0 * np.pi * var)) + np.sum((X - mean) ** 2
124 / var, axis=1))
125
126     def _e_step(self, X: np.ndarray, w: np.ndarray, m: np.ndarray, v: np.ndarray) -> Tuple[np.ndarray, float]:
127         n = X.shape[0]
128         logp = np.empty((n, self.C), dtype=np.float64)
129         for c in range(self.C):
130             logp[:, c] = np.log(w[c] + 1e-300) + self._log_gauss_diag(X, m[c], v[c])
131         log_norm = logsumexp(logp, axis=1)
132         return logp - log_norm[:, None], float(np.mean(log_norm))
133
134     def _m_step(self, X: np.ndarray, log_resp: np.ndarray) -> Tuple[np.ndarray, np.ndarray, np.ndarray]:
135         resp = np.exp(log_resp)
136         n, d = X.shape
137         Nk = resp.sum(axis=0) + 1e-12
138         w = Nk / n
139         m = (resp.T @ X) / Nk[:, None]
140         v = np.empty((self.C, d), dtype=np.float64)
141         for c in range(self.C):
142             diff = X - m[c]
143             v[c] = (resp[:, c:c+1] * (diff ** 2)).sum(axis=0) / Nk[c]
144             v[c] = np.maximum(v[c] + self.reg, self.min_var)
145
146     def _fit_once(self, X: np.ndarray, rng: np.random.Generator) -> Tuple[np.ndarray, np.ndarray, np.ndarray, float]:
147         n, d = X.shape
148         idx = rng.choice(n, size=self.C, replace=False)

```

```

149     m = X[idx].copy() + rng.normal(scale=1e-3, size=(self.C, d))
150     base_var = np.maximum(np.var(X, axis=0) + self.reg, self.min_var)
151     v = np.tile(base_var[None, :], (self.C, 1))
152     w = np.full((self.C,), 1.0 / self.C, dtype=np.float64)
153
154     prev = -np.inf
155     best = -np.inf
156     for _ in range(self.n_iter):
157         log_resp, lb = self._e_step(X, w, m, v)
158         if lb > best:
159             best = lb
160         if np.isfinite(prev) and abs(lb - prev) < self.tol:
161             break
162         prev = lb
163         w, m, v = self._m_step(X, log_resp)
164     return w, m, v, best
165
166 def fit_best_of_n(self, X: np.ndarray, n_init: int, seed: int) -> "GMMDiagonal":
167     rngg = np.random.default_rng(seed)
168     best_lb = -np.inf
169     best_params = None
170     for _ in range(max(1, int(n_init))):
171         rng = np.random.default_rng(int(rngg.integers(0, 2**31 - 1)))
172         w, m, v, lb = self._fit_once(X, rng)
173         if lb > best_lb:
174             best_lb = lb
175             best_params = (w, m, v)
176     self.w, self.m, self.v = best_params
177     return self
178
179 def score_samples(self, X: np.ndarray) -> np.ndarray:
180     n = X.shape[0]
181     logp = np.empty((n, self.C), dtype=np.float64)
182     for c in range(self.C):
183         logp[:, c] = np.log(self.w[c] + 1e-300) + self._log_gauss_diag(X, self.m[c], self.v[c])
184     return logsumexp(logp, axis=1)
185
186
187 def classify(fg: GMMDiagonal, bg: GMMDiagonal, X: np.ndarray, p_fg: float, p_bg: float) -> np.ndarray:
188     ll_fg = fg.score_samples(X) + np.log(p_fg + 1e-300)
189     ll_bg = bg.score_samples(X) + np.log(p_bg + 1e-300)
190     return (ll_fg > ll_bg).astype(np.uint8)
191
192
193 def poe(pred: np.ndarray, gt: np.ndarray) -> float:
194     return float(np.mean(pred.astype(np.uint8) != gt.astype(np.uint8)))
195
196
197 def solve(
198     p: Paths,
199     dims: List[int],
200     c_list: List[int],
201     n_runs: int,
202     seed: int,
203     max_em_iter: int,
204     tol: float,

```

```

205     reg: float,
206     min_var: float,
207     n_init_b: int,
208 ) -> None:
209     os.makedirs(p.out_dir, exist_ok=True)
210
211     print("[INFO] Loading data...")
212     train_fg, train_bg = load_training(p.train_mat)
213     img, mask = load_img_mask(p.img, p.mask)
214
215     prior_fg = train_fg.shape[0] / (train_fg.shape[0] + train_bg.shape[0])
216     prior_bg = 1.0 - prior_fg
217     print(f"[INFO] Training priors: FG={prior_fg:.6f}, BG={prior_bg:.6f}")
218
219     order = zigzag_order(p.zigzag)
220     print("[INFO] Extracting test image features...")
221     Xtest, h2, w2 = dct_features_valid(img, order, dc_scale=0.5)
222     ytest = mask_valid(mask, h2, w2)
223     print(f"[INFO] eval size={h2}x{w2} (N={h2*w2})")
224
225     feat_rank = alpha_ranking_from_subsets(p.subsets_mat)
226     print("[INFO] Feature ranking source: computed alpha from subsets D4_FG/D4_BG")
227 )
228
229 def cols_for_d(d: int) -> np.ndarray:
230     return feat_rank[:d]
231
232 rng_global = np.random.default_rng(seed)
233
234     print("\n[INFO] --- Part (a): Initialization sensitivity (C=8) ---")
235     C_a = 8
236
237     fg_models_by_d: List[List[GMMDiagonal]] = []
238     bg_models_by_d: List[List[GMMDiagonal]] = []
239
240     for d in dims:
241         cols = cols_for_d(d)
242         Xfg_d = train_fg[:, cols]
243         Xbg_d = train_bg[:, cols]
244
245         fg_models = []
246         bg_models = []
247
248         for _ in range(n_runs):
249             s = int(rng_global.integers(0, 2**31 - 1))
250             fg_models.append(GMMDiagonal(C_a, max_em_iter, tol, reg, min_var).
251             fit_best_of_n(Xfg_d, 1, s))
252
253             for _ in range(n_runs):
254                 s = int(rng_global.integers(0, 2**31 - 1))
255                 bg_models.append(GMMDiagonal(C_a, max_em_iter, tol, reg, min_var).
256                 fit_best_of_n(Xbg_d, 1, s))
257
258             fg_models_by_d.append(fg_models)
259             bg_models_by_d.append(bg_models)
260
261     err_bgfix = [[[0.0 for _ in dims] for _ in range(n_runs)] for _ in range(
262     n_runs)]
263     for k, d in enumerate(dims):

```

```

260     cols = cols_for_d(d)
261     Xte_d = Xtest[:, cols]
262     for j in range(n_runs):
263         for i in range(n_runs):
264             pred = classify(fg_models_by_d[k][i], bg_models_by_d[k][j], Xte_d,
265             prior_fg, prior_bg)
266             err_bfgfix[j][i][k] = poe(pred, ytest)
267
268     for k, d in enumerate(dims):
269         all_pairs = np.array([err_bfgfix[j][i][k] for j in range(n_runs) for i in
270             range(n_runs)], dtype=np.float64)
271         print(f"[INFO] d={d:2d}: mean={all_pairs.mean():.6f} (min={all_pairs.min():
272             .6f}, max={all_pairs.max():.6f})")
273
274     fig, axes = plt.subplots(2, 3, figsize=(14, 8), constrained_layout=True)
275     axes = axes.flatten()
276     for j in range(n_runs):
277         ax = axes[j]
278         for i in range(n_runs):
279             ax.plot(dims, err_bfgfix[j][i], marker="o", linewidth=1)
280             ax.set_title(f"PE vs d (BG #{j+1} fixed)")
281             ax.set_xlabel("Feature number (d)")
282             ax.set_ylabel("Probability of Error")
283             ax.grid(True)
284     axes[-1].axis("off")
285
286     out_a = os.path.join(p.out_dir, "prob6_a_initialization.png")
287     fig.suptitle("Part (a): Probability of Error vs Feature number (BG fixed; 5 FG
288     curves)", fontsize=14)
289     fig.savefig(out_a, dpi=150, bbox_inches="tight")
290     plt.close(fig)
291     print(f"[INFO] Saved: {out_a}")
292
293     print("\n[INFO] --- Part (b): Vary number of components C ---")
294     plt.figure(figsize=(12, 8))
295
296     for C in c_list:
297         errors_c = []
298         for d in dims:
299             cols = cols_for_d(d)
300             Xfg_d = train_fg[:, cols]
301             Xbg_d = train_bg[:, cols]
302             Xte_d = Xtest[:, cols]
303
304             s_fg = int(rng_global.integers(0, 2**31 - 1))
305             s_bg = int(rng_global.integers(0, 2**31 - 1))
306
307             fg_model = GMMDiagonal(C, max_em_iter, tol, reg, min_var).
308             fit_best_of_n(Xfg_d, n_init_b, s_fg)
309             bg_model = GMMDiagonal(C, max_em_iter, tol, reg, min_var).
310             fit_best_of_n(Xbg_d, n_init_b, s_bg)
311
312             pred = classify(fg_model, bg_model, Xte_d, prior_fg, prior_bg)
313             errors_c.append(poe(pred, ytest))
314
315             plt.plot(dims, errors_c, marker="o", label=f"C={C}")
316             print(f"[INFO] C={C:2d}: errors = {[f'{e:.4f}' for e in errors_c]}}")

```

```

312     plt.title("Part (b): Probability of Error vs Feature number (varying
313     components)")
314     plt.xlabel("Feature number (d)")
315     plt.ylabel("Probability of Error")
316     plt.grid(True)
317     plt.legend()
318
319     out_b = os.path.join(p.out_dir, "prob6_b_components.png")
320     plt.savefig(out_b, dpi=150, bbox_inches="tight")
321     plt.close()
322     print(f"[INFO] Saved: {out_b}")
323
324
325
326 def main() -> None:
327     ap = argparse.ArgumentParser()
328     ap.add_argument("--data-dir", default="data")
329     ap.add_argument("--output-dir", default="hw4/output")
330     ap.add_argument("--dims", default=",".join(map(str, DIMS_DEFAULT)))
331     ap.add_argument("--c-list", default=",".join(map(str, C_LIST_DEFAULT)))
332     ap.add_argument("--n-runs", type=int, default=5)
333     ap.add_argument("--seed", type=int, default=0)
334     ap.add_argument("--max-em-iter", type=int, default=200)
335     ap.add_argument("--tol", type=float, default=1e-4)
336     ap.add_argument("--reg", type=float, default=1e-6)
337     ap.add_argument("--min-var", type=float, default=1e-6)
338     ap.add_argument("--n-init-b", type=int, default=5)
339     args = ap.parse_args()
340
341     p = Paths(
342         train_mat=os.path.join(args.data_dir, TRAIN_MAT),
343         subsets_mat=os.path.join(args.data_dir, SUBSETS_MAT),
344         img=os.path.join(args.data_dir, IMG),
345         mask=os.path.join(args.data_dir, MASK),
346         zigzag=os.path.join(args.data_dir, ZIGZAG),
347         out_dir=args.output_dir,
348     )
349
350     solve(
351         p=p,
352         dims=parse_list_int(args.dims),
353         c_list=parse_list_int(args.c_list),
354         n_runs=args.n_runs,
355         seed=args.seed,
356         max_em_iter=args.max_em_iter,
357         tol=args.tol,
358         reg=args.reg,
359         min_var=args.min_var,
360         n_init_b=args.n_init_b,
361     )
362
363
364 if __name__ == "__main__":
365     main()

```

Listing 1: Python code for HW4